


CS299 Lecture Notes 1

Supervised Learning



- input $x^{(i)}$ (also called input features)
- output $y^{(i)}$ (or target variable)
- A pair $(x^{(i)}, y^{(i)})$ is called a training example
- Training Set $\{(x^{(i)}, y^{(i)}) : i = 1, \dots, m\}$
- Input and output space $X = Y (= \mathbb{R}$ in our case)
- we wish to learn a function $h: X \rightarrow Y$ so that h is a good predictor

Linear Regression

$$h(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

- by convention, $x_0 = 1$
- Given a training set, how do we pick h ? i.e., the params θ
- we can try to make $h(x)$ as close as possible to y , at least for the training examples that we have. we define the cost function

$$\begin{aligned} J(\theta) &= \frac{1}{2} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2 \\ &= \frac{1}{2} \|X\theta - y\|_2^2 \end{aligned}$$

$$\text{where } X = \begin{pmatrix} - & x^{(1)T} \\ & \vdots \\ - & x^{(m)T} \end{pmatrix}$$

LMS Algorithm

- we can minimize $J(\theta)$ by gradient descent

m training examples
 n features ($n+1$ if including the intercept x_0)

- start w/ some θ , and repeatedly update by:

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

for a single example:

$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{\partial}{\partial \theta_j} \frac{1}{2} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2 = \boxed{(h_\theta(x) - y) x_j}$$

for the whole training set: $\frac{\partial}{\partial \theta} J(\theta) = \frac{\partial}{\partial \theta} \frac{1}{2} \|A\theta - y\|^2$

$$= 2 \cdot \frac{1}{2} A^T (A\theta - y)$$

$$= A^T A \theta - A^T y \stackrel{!}{=} 0$$

$$\boxed{\theta = (A^T A)^{-1} A^T y}$$

- this is called the LMS update rule, or Widrow-Hoff.
- update is proportional to the error term $(y^{(i)} - h_\theta(x^{(i)}))$
- Batch gradient descent: use whole dataset ~~all~~ for each update
- Stochastic gradient descent: update parameters w.r.t a subset of examples.

- Matrix derivatives

$$\nabla_A \text{tr}(AB) = B^T \quad (A_{n \times m}, B_{m \times n})$$

$$\nabla_{A^T} f(A) = (\nabla_A f(A))^T$$

$$\nabla_A \text{tr} ABA^T C = CAB + C^T A B^T$$

$$\nabla_A |A| = |A| (A^{-1})^T \quad |A| - \text{determinant}$$

Normal Equation

$$X = \begin{bmatrix} \text{---} (x^{(1)})^T \text{---} \\ \vdots \\ \text{---} (x^{(m)})^T \text{---} \end{bmatrix} \in \mathbb{R}^{m \times n}$$

$$y = \begin{bmatrix} y^{(1)} \\ \vdots \\ y^{(m)} \end{bmatrix}$$

$$J(\theta) = \frac{1}{2} (X\theta - y)^T (X\theta - y) = \frac{1}{2} \|X\theta - y\|^2$$

probabilistic interpretation

- What are the assumptions under which linear regression is a reasonable choice?

$$y^{(i)} = \theta^T x^{(i)} + \epsilon^{(i)}$$

Where $\epsilon^{(i)}$ is an error term that captures unmodeled effects (i.e., missing features), or random noise.

- Let us further assume that the $\epsilon^{(i)}$ are distributed IID, according to a Gaussian dist. with $(0, \sigma^2)$. We can write

$$\epsilon^{(i)} \sim N(0, \sigma^2)$$

and the density is given by

$$p(\epsilon^{(i)}) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(\epsilon^{(i)})^2}{2\sigma^2}\right)$$

$$\Rightarrow p(y^{(i)} | x^{(i)}; \theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right)$$

- Denote

$$L(\theta) = L(\theta; X, y) = p(y | X; \theta)$$

that is, given a certain dataset X , we can calculate the likelihood of any prediction y .

- by the IID assumption

$$L(\theta) = \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta)$$

$$= \prod \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right)$$

- We maximize L , or $\ell = \log(L)$

$$\begin{aligned}
 \ell(\theta) &= \log L(\theta) \\
 &= \log \prod_{i=1}^m \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\
 &= \sum_{i=1}^m \log \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\
 &= \sum_{i=1}^m \log \frac{1}{\sqrt{2\pi}\sigma} + \log \exp\left(-\frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}\right) \\
 &= m \log \frac{1}{\sqrt{2\pi}\sigma} - \sum_{i=1}^m \frac{(y^{(i)} - \theta^T x^{(i)})^2}{2\sigma^2}
 \end{aligned}$$

• So, to maximize $\ell(\theta)$, we can have to minimize $\sum_{i=1}^m (y^{(i)} - \theta^T x^{(i)})^2$

• So, under the assumption $y^{(i)} = x^{(i)} + \epsilon^{(i)}$, doing least squares regression corresponds to finding the maximum likelihood estimate of θ .

Locally weighted Linear Regression

- underfitting: the data has structure not captured by the model.
- overfitting: the model has structure not inherent to the data.

• fit θ to minimize $\sum_i w^{(i)} (y^{(i)} - \theta^T x^{(i)})^2$
output $\theta^T x$

where $w^{(i)} = \exp\left(-\frac{(x^{(i)} - x)^2}{2\tau^2}\right)$

which gives higher weight to pts close to the query pt.


- Parametric
- Non-parametric: amount of stuff we need to keep track of

Logistic Regression

- Suppose now that $y^{(i)} \in \{0, 1\}$
- No longer makes sense to use linear regression
- Instead, we will choose

$$h_{\theta}(x) = g(\theta^T x) = \frac{1}{1 + e^{-\theta^T x}}$$

$$g(z) = \frac{1}{1 + e^{-z}}$$



(called the sigmoid or logistic function)

- $g'(z) = g(z)(1 - g(z))$
- The starting assumption (probabilistic) is that

$$p(y=1 | x; \theta) = h_{\theta}(x)$$

$$p(y=0 | x; \theta) = 1 - h_{\theta}(x)$$

which can be written more compactly as

$$p(y | x; \theta) = (h_{\theta}(x))^y (1 - h_{\theta}(x))^{1-y}$$

- The likelihood is now

$$L(\theta) = \prod_{i=1}^m p(y^{(i)} | x^{(i)}; \theta)$$

$$= p(\vec{y} | X; \theta)$$

$$= \prod_{i=1}^m (y^{(i)} | x^{(i)}; \theta)$$

$$= \prod (h_{\theta}(x^{(i)})^{y^{(i)}} (1 - h_{\theta}(x^{(i)}))^{1-y^{(i)}})$$

~~max~~ Taking the log likelihood and maximizing it with SGD, we get the update step:

$$\theta_J = \theta_J + \alpha (y^{(i)} - h_{\theta}(x^{(i)})) x_J^{(i)}$$

which is suspiciously similar to LMS ... (spoiler: GLM)

Digression: The perceptron Learning Algorithm

- Consider modifying the sigmoid function to "force" it to output 0 or 1.

$$g(z) = \begin{cases} 1 & \text{if } z \geq 0 \\ 0 & \text{if } z < 0 \end{cases}$$

- Using $h_{\theta}(x) = g(\theta^T x)$ and the update rule

$$\theta_j := \theta_j + \alpha (\gamma^{(i)} - h_{\theta}(x^{(i)})) x_j^{(i)}$$


We get what is called the perceptron algorithm.

- In the 60s, this was argued to be a rough model of how a single neuron in the brain worked.
- Good starting pt for theory analysis
- Hard to derive as a max likelihood sort of estimation alg

Another Algorithm for Maximizing $\ell(\theta)$

- Consider logistic regression again.
- Newton's method uses a linear approx of a function $f: \mathbb{R} \rightarrow \mathbb{R}$ to find a θ s.t. $f(\theta) = 0$

$$\theta := \theta - \frac{f(\theta)}{f'(\theta)}$$



$$\begin{aligned} \frac{f(\theta)}{x} &= f'(\theta) \\ x &= \frac{f(\theta)}{f'(\theta)} \end{aligned}$$

- What if we want to use it to maximize $\ell(\theta)$?
We can try to find $\ell'(\theta) = 0$ using Newton's method:

$$\theta := \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$$

- The generalization of Newton's method to the multidimensional setting is given by

$$\boxed{\theta := \theta - H^{-1} \nabla_{\theta} \ell(\theta)}$$

Where H is the Hessian, $H_{ij} = \frac{\partial^2 \ell(\theta)}{\partial \theta_i \partial \theta_j}$

- Newton's method usually takes fewer iterations to converge, but each iteration is more expensive because of H^{-1} .

Generalized Linear Models

- Both the regression $(y|x; \theta \sim N(\mu, \sigma^2))$ and the classification $(y|x; \theta \sim \text{Bernoulli}(\phi))$ examples are a special case of a GLM.

The Exponential Family

- We start by defining the exponential family of distributions

$$p(y; \eta) = b(y) \exp(\eta^T T(y) - a(\eta))$$

η - natural parameter (also called canonical param.)

$T(y)$ - sufficient statistic (often $T(y) = y$)

$a(\eta)$ - log partition function. $e^{-a(\eta)}$ plays the role of normalization, making sure the dist. sums/integrates to 1.
(did he mean $b(y)$?)

T, a, b defines a family of distributions that is parametrized by η .

- The Bernoulli and Gaussian distr. are exp. family dist.
- Bernoulli:

$$\begin{aligned} p(y; \phi) &= \phi^y (1-\phi)^{1-y} \\ &= \exp(\log \phi^y (1-\phi)^{1-y}) \\ &= \exp(y \log \phi + (1-y) \log (1-\phi)) \\ &= \exp(\underbrace{y}_{T(y)} \underbrace{\log \frac{\phi}{1-\phi}}_{\eta} + \underbrace{\log (1-\phi)}_a) \end{aligned}$$

$$T(y) = y$$

$$a(\eta) = -\log(1-\phi)$$

$$\boxed{\phi = \frac{1}{1+e^{-\eta}}}$$

$$= -\log\left(1 - \frac{1}{1+e^{-\eta}}\right) = -\log\left(\frac{e^{-\eta}}{1+e^{-\eta}}\right)$$

$$= \log\left(\frac{1+e^{-\eta}}{e^{-\eta}}\right) = \log(1+e^{\eta})$$

$$b(y) = 1$$

- Gaussian: for convenience, choose $\sigma^2 = 1$ (it had no affect)

$$p(y; \mu) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}(y-\mu)^2\right)$$

$$= \underbrace{\frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}y^2\right)}_b \exp\left(\underbrace{\mu y}_{\eta \tau} - \underbrace{\frac{1}{2}\mu^2}_a\right)$$

$$\eta = \mu$$

$$\tau(y) = y$$

$$a(\eta) = \frac{1}{2}\mu^2 = \frac{1}{2}\eta^2$$

$$b(y) = \frac{1}{\sqrt{2\pi}} \exp\left(-\frac{1}{2}y^2\right)$$

Constructing GLMs

- We want to predict y given x .
- To derive a GLM, we make the following assumptions:
 1. $y|x; \theta \sim \text{Exponential Family}(\eta)$
 2. Given x , ^{our} goal is to predict the expected value of $T(y)$. ^(given x) Usually, $T(y) = y$, which means we want h to satisfy $h_\theta(x) = E[y|x]$. For example, in logistic regression $h_\theta(x) = p(y=1|x) = E[y|x; \theta]$.
 3. The natural param. and x are related by $\eta = \theta^T x$ (this is more of a design choice)
- These assumptions allow us to derive a very elegant class of learning algs. In particular, we can derive logistic regression and ordinary least squares.

Ordinary Least Squares

- y is continuous
- We model $y|x; \theta \sim N(\mu, \sigma^2)$
- So, we let the Exponential Family(η) be the normal dist. as we saw, $\eta = \mu$. So, we have:

$$h_\theta(x) = E[y|x; \theta] \quad (\text{assumption 2})$$

$$= \mu$$

$$= \eta = \theta^T x \quad (\text{assumption 3})$$

Logistic Regression

- $y \in \{0, 1\}$
- $\phi = \frac{1}{1+e^{-\eta}}$
- $y|x; \theta \sim \text{Bernoulli}(\phi) \Rightarrow E[y|x; \theta] = \phi$
- we get

$$\begin{aligned} h_{\theta}(x) &= E[y|x; \theta] \\ &= \phi \\ &= \frac{1}{1+e^{-\eta}} \\ &= \frac{1}{1+e^{-\theta^T x}} \end{aligned}$$

- So, assuming a Bernoulli dist. and an exponential and the GLM assumptions together give rise to the sigmoid hypothesis.
- canonical response function $g(\eta) = E[T(y); \eta]$
- canonical inverse function g^{-1}

Softmax Regression

- $y \in \{1, 2, \dots, K\}$
- probability of each class is parametrized by $\phi_1, \dots, \phi_{K-1}$
- $\phi_K = 1 - \sum_{i=1}^{K-1} \phi_i$ (fully specified by $\phi_1, \dots, \phi_{K-1}$)
- $P(y=i | \phi) = \phi_i$
- To express as an exponential family, define $T(y) \in \mathbb{R}^{K-1}$

$$T(1) = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \dots, T(K-1) = \begin{bmatrix} 0 \\ \vdots \\ 1 \\ 0 \end{bmatrix}, T(K) = 0$$

- Unlike prev examples, $T(y)$ is a vector.
- $(T(y))_i = 1 \{y=i\}$
- $E[(T(y))_i] = P(y=i) = \phi_i$
- The multinomial is a member of the exponential family. we have

$$\begin{aligned}
p(y; \phi) &= \phi_1^{1\{y=1\}} \phi_2^{1\{y=2\}} \dots \phi_k^{1\{y=k\}} \\
&= \phi_1^{1\{y=1\}} \phi_2^{1\{y=2\}} \dots \phi_k^{1 - \sum_{i=1}^{k-1} 1\{y=i\}} \\
&= \phi_1^{(T(y))_1} \phi_2^{(T(y))_2} \dots \phi_k^{1 - \sum_{i=1}^{k-1} (T(y))_i} \\
&= \exp((T(y))_1 \log \phi_1 + (T(y))_2 \log \phi_2 + \dots + (1 - \sum_{i=1}^{k-1} (T(y))_i) \log \phi_k) \\
&= \exp((T(y))_1 \log \frac{\phi_1}{\phi_k} + \dots + (T(y))_{k-1} \log \frac{\phi_{k-1}}{\phi_k} + \log \phi_k) \\
&= b(y) \exp(\eta^T T(y) - a(\eta))
\end{aligned}$$

where

$$\eta = \begin{bmatrix} \log(\phi_1 / \phi_k) \\ \log(\phi_2 / \phi_k) \\ \vdots \\ \log(\phi_{k-1} / \phi_k) \end{bmatrix} \quad (\eta_k = 0 \text{ for convenience})$$

$$a(\eta) = -\log \phi_k$$

$$b(y) = 1$$

• This completes our formulation of the multinomial as an exponential family distribution.

• The link function is given by $\eta_i = \log \frac{\phi_i}{\phi_k}$

• To invert the link func:

$$e^{\eta_i} = \frac{\phi_i}{\phi_k}$$

$$\phi_k \sum_{i=1}^k e^{\eta_i} = \sum_{i=1}^k \phi_i = 1$$

$$\Rightarrow \boxed{\phi_k = \frac{1}{\sum_{i=1}^k e^{\eta_i}}}$$

$$\Rightarrow \boxed{\phi_i = \frac{e^{\eta_i}}{\sum_{j=1}^k e^{\eta_j}}}$$

softmax function

• by assumption 3, $\eta_i = \theta_i^T x$ for $i=1, \dots, k-1$

where $\theta_1, \dots, \theta_{k-1} \in \mathbb{R}^{n+1}$ are the params of our model

• for convenience, we also define $\theta_k = 0$, so that $\eta_k = \theta_k^T x = 0$

- Hence, our model assumes that the conditional dist. of y given x is given by

$$\begin{aligned} p(y=i | x; \theta) &= \phi_i \\ &= \frac{e^{\eta_i}}{\sum_{i=1}^K e^{\eta_i}} \\ &= \frac{\exp(\theta_i^T x)}{\sum_{j=1}^K \exp(\theta_j^T x)} \end{aligned}$$

- Our hypothesis will output

$$h_\theta(x) = E[T(y) | x, \theta] = E \begin{bmatrix} 1 \text{ if } y=1 \\ 1 \text{ if } y=2 \\ \vdots \\ 1 \text{ if } y=K-1 \end{bmatrix} \bigg| x; \theta = \begin{bmatrix} \phi_1 \\ \phi_2 \\ \vdots \\ \phi_{K-1} \end{bmatrix} = \begin{bmatrix} \dots \end{bmatrix}$$

- Finally, we fit the parameters by maximizing

$$\begin{aligned} \ell(\theta) &= \sum_{i=1}^m \log p(y^{(i)} | x^{(i)}; \theta) \\ &= \sum_{i=1}^m \log \prod_{k=1}^K \left(\frac{\exp(\theta_k^T x^{(i)})}{\sum_{j=1}^K \exp(\theta_j^T x^{(i)})} \right)^{1 \text{ if } y^{(i)}=k} \end{aligned}$$

(using gradient ascent or Newton's method)